Keywords: explorative modeling • generative pretraining • mode commitment • sample efficiency

Explorative Modeling: Unlocking a Third Pretraining Axis and End-to-End Generation

Training Generative Models to Commit to Modes by Sampling K Candidates and Updating Only the Best—Yielding 6× Sample Efficiency at Near-SOTA Quality

By Alexi Gladstone, Heng Ji, Yilun Du (University of Illinois Urbana-Champaign; MIT)

arXiv:2607.27372 • July 29, 2026

Generative model training has two well-known scaling axes: parameters and data. A new paper proposes a third—exploration—by sampling K candidate outputs per training example and updating only on the best match. **Explorative Models (XMs)** reach 6.2× sample efficiency, 4.1× FLOP efficiency, and 47% better parameter efficiency compared to standard training, with a near-SOTA FID of 1.43 on ImageNet 256×256 without guidance.

AT A GLANCE

Problem: Standard MLE training averages over modes of the data distribution, producing blurry or incoherent generations.

Why it matters: Mode averaging is a fundamental inefficiency in generative training—fixing it unlocks major sample and FLOP savings.

Method: Sample K candidates per training example; update only on the best-matching one, so gradients reinforce mode-committing outputs.

Result: 6.2× sample efficiency, 4.1× FLOP efficiency, 47% parameter savings, FID 1.43 on ImageNet.

Evidence: Gains are monotone in K across images, video, and language; advantage holds in large-scale pretraining.

The Big Picture

The deep learning revolution was built on end-to-end training: instead of decomposing a problem into hand-designed stages, a single model is trained end-to-end by gradient descent on a task loss. Generative modeling has remained a partial exception. While diffusion models, autoregressive models, and flow models have made remarkable progress, the training loop itself has remained standard: minimize the expected loss over all data examples, averaging over all the modes in the distribution.

This paper argues that the averaging implicit in maximum likelihood estimation is a fundamental inefficiency. When a training example has multiple plausible completions, the gradient pushes the model to cover all of them simultaneously—producing predictions that blur across modes rather than committing to one. The proposed fix, Explorative Modeling (XM), factors the training loop to commit to modes: sample K candidates, pick the best, update on that. Exploration is treated as a quantitative scaling axis alongside parameters and data.

Why Existing Approaches Fall Short

Maximum likelihood training minimizes

$$\mathcal{L}_{\mathrm{MLE}}(\theta) = -\mathbb{E}_{(x,c)\sim\mathcal{D}}\left[\log p_\theta(x \mid c)\right]$$

This encourages the model to assign probability mass to every mode of the data distribution. In practice, this causes mode averaging: the model hedges rather than commits, producing blurry images, inconsistent video frames, or generic text.

Existing remedies address the symptom at inference time (guidance, best-of-N sampling) but leave the training objective unchanged. This means mode-averaging gradients are incorporated into model weights, requiring extra inference compute to overcome them. XM addresses the root cause at training time.

Core Mechanism

At each training step, instead of computing a loss from one generation matched to one data example, XM:

1. Samples K candidate generations $\hat{x}_1, \dots, \hat{x}_K \sim p_\theta(\cdot \mid c)$.
2. Selects the best-matching candidate: $\hat{x}^* = \arg\max_k \operatorname{sim}(\hat{x}_k, x)$.
3. Computes the training gradient only on $\hat{x}^*$.

The key invariant: gradients always reinforce mode-committing outputs. The model is never penalized for failing to cover modes it did not select. For $K = 1$, XM reduces to standard MLE; increasing K monotonically improves training quality.

Technical Formulation

Let $p_\theta(x \mid c)$ be the generative model. The XM objective is:

$$\mathcal{L}_{\mathrm{XM}}(\theta) = -\mathbb{E}_{(x,c)\sim\mathcal{D}}\left[\log p_\theta(\hat{x}^* \mid c)\right]$$

where

$$\hat{x}^* = \arg\max_{\hat{x}_k \in \{\hat{x}_1,\dots,\hat{x}_K\}} \operatorname{sim}(\hat{x}_k,\, x).$$

The similarity function $\mathrm{sim}(\cdot,\cdot)$ is task-dependent: perceptual distance for images, video quality metrics for video, and ROUGE/token-overlap for language. The gradient is:

$$\nabla_\theta \mathcal{L}_{\mathrm{XM}} = -\nabla_\theta \log p_\theta(\hat{x}^* \mid c)$$

—identical in form to MLE but computed only on the best candidate.

The authors show that $\mathcal{L}_{\mathrm{XM}}$ upper-bounds the negative log-likelihood of the best-of-K sample. As $K \to \infty$, the objective converges to the entropy of the highest-probability mode, incentivizing mode commitment rather than mode coverage.

Learning or Inference Procedure

Training is identical to standard generative model training except for the sampling step. No changes are needed to the architecture, optimizer, or data pipeline. The per-step compute cost increases by a factor of approximately K for forward passes (to generate K candidates), but only one backward pass is performed, so the memory overhead is modest.

At inference, the model is used identically to a standard generative model: single forward pass, no additional ranking or sampling. Exploration is baked into the parameters at training time, not deferred to test time.

What the Guarantee Says

For $K = 1$, $\mathcal{L}_{\mathrm{XM}} = \mathcal{L}_{\mathrm{MLE}}$. For $K > 1$, $\mathcal{L}_{\mathrm{XM}} \leq \mathcal{L}_{\mathrm{MLE}}$—the XM objective is never worse, and strictly better when the model has non-zero variance over candidates. The gap grows with the diversity of the model's distribution: early in training, when the model is uncertain, exploration provides the largest benefit.

Experimental Findings

ImageNet 256×256 generation (FID, lower is better):

Model	FID	Guidance
Standard MLE baseline	2.21	No
XM ($K = 8$)	**1.43**	No
XM ($K = 8$) + CFG	**1.41**	Yes
SOTA diffusion model	1.40	Yes

Efficiency at matched FID (FID = 2.21):

Axis	XM vs. Standard
Sample efficiency	6.2×
FLOP efficiency	4.1×
Parameter efficiency	47% fewer params

Scaling of exploration (K): Increasing $K \in \{1, 2, 4, 8, 16, 32\}$ monotonically reduces FID, with approximately log-linear improvement.

Video generation: XM reduces FVD by 21% relative to the MLE baseline at matched compute on a text-to-video benchmark.

Language generation: XM improves the diversity-quality Pareto frontier by 15% MAUVE on a causal LM fine-tuning task.

Ablations and Interpretation

Similarity metric: Replacing perceptual loss with L2 pixel distance degrades FID by 0.6, confirming that the quality of the match function determines how well candidates proxy for mode quality.

K vs. additional training steps: At a fixed total FLOP budget, $K = 4$ with more steps outperforms $K = 1$ with fewer steps, confirming that exploration is more efficient than additional data processing when the model's own distribution spans the relevant modes.

Connection to best-of-N: Best-of-N at inference requires N forward passes per generation and does not affect model weights. XM amortizes the exploration cost into training: the inference model inherits mode-committing behavior without any per-generation overhead.

End-to-end generation: XM stabilizes joint training of encoder-decoder architectures with shared parameters across intermediate and final outputs, because mode-committing gradients prevent the latent space from collapsing toward averaged representations.

Reference

Alexi Gladstone, Heng Ji, Yilun Du. **Explorative Modeling: Unlocking a Third Pretraining Axis and End-to-End Generation**. arXiv:2607.27372, July 29, 2026. https://arxiv.org/abs/2607.27372