Keywords: multimodal pretraining • knowledge transfer • modality synergy • vision laziness

Towards Physics of Multimodal Pretraining: Knowledge Flow, Modality Synergy, Early Unification, and Recipes

Systematic Empirical Study Reveals Asymmetric Knowledge Flow, Data-Driven Synergy, and a Vision Laziness Phenomenon That Together Explain Why Early Joint Training Dominates

By Junlin Han, Shengbang Tong, David Fan, Minghao Chen, Philip Torr, Filippos Kokkinos, Mike Lewis

arXiv:2608.05000 • Preprint, August 2026

Training a single model to understand and generate across modalities has become the dominant paradigm in multimodal AI, yet the mechanisms by which modalities interact during joint pretraining remain poorly understood. This paper presents a systematic empirical investigation — a “physics” of multimodal pretraining — identifying four mechanistic principles that explain why certain design choices succeed: asymmetric knowledge flow, data-complexity-driven synergy, early unification superiority, and compute-efficient recipes that deliver strong models at 5% of naive training cost.

AT A GLANCE

Problem: Multimodal pretraining design is largely empirical; the mechanisms behind successful choices are unknown.

Why it matters: Principled design could drastically reduce the compute required to train capable multimodal models.

Method: Controlled ablations over modality combinations, training order, and architecture choices; validated by training 13.5B MoE models on 2T tokens.

Result: Four mechanistic principles identified; efficient recipes achieve competitive quality at 5% of naive joint-training compute.

Evidence: Cross-modal transfer measured via held-out evaluation; vision laziness confirmed by representation analysis.

The Big Picture

Modern multimodal models are trained to process and generate across language, visual understanding (image captioning, visual question answering), and visual generation (text-to-image synthesis). The standard practice is joint pretraining on a mixture of all modalities simultaneously. This works, but the design space is vast: which modalities to combine, in what order, with what architecture, and for how long on each task. Each choice is typically resolved through expensive ablations with no theoretical grounding.

This paper asks: can we derive principled design rules from the mechanisms of multimodal training? By systematically varying one factor at a time and measuring the impact on all modalities, the authors identify four mechanistic principles that collectively constitute a “physics” of multimodal pretraining.

Why Existing Approaches Fall Short

The most common baselines for multimodal model design are:

Naive joint training: Mix all modality data from the start in fixed ratios. Works, but expensive and provides no insight into which choices matter.

Sequential pretraining: Pretrain on language, then adapt to vision via fine-tuning (e.g., LLaVA-style). Efficient but caps multimodal integration quality because the language representations are already frozen when vision is introduced.

Late alignment: Train unimodal encoders separately, then align them via a projection layer. Limited by the quality of the alignment stage and the incompatibility of separately-optimized representation spaces.

None of these approaches is informed by an understanding of *how* modalities interact. The four principles below provide that understanding.

Core Mechanism

Principle 1 — Asymmetric Knowledge Flow. Language knowledge transfers readily into visual understanding (via shared text-image pairs), but visual generation benefits less directly from language pretraining. The asymmetry means language pretraining is a high-value starting point for understanding tasks but should not be allowed to dominate the training objective for generation tasks.

Principle 2 — Data Complexity Drives Synergy. On simple datasets (small vocabulary, low intra-class variability), training on multiple modalities simultaneously induces modality competition: one modality dominates gradient updates and the others underfit. On complex, diverse datasets, the modalities are synergistic: each modality provides complementary learning signal that improves performance on the others. Architecturally, synergy is promoted by shared attention and normalization layers with modality-specific feed-forward networks.

Principle 3 — Early Unification and Vision Laziness. Unifying modalities from the start of training outperforms sequential or late alignment. The paper identifies the mechanism: *vision laziness*. When visual input is introduced after strong language representations have formed, the model has already developed robust language priors. It then relies on those priors rather than learning genuinely visual representations. The visual encoder receives gradient signals that push it to produce language-compatible features rather than visually grounded ones. Early unification prevents this by forcing visual and language representations to co-evolve from the first gradient step.

Principle 4 — Efficient Recipes. Combining the first three principles, the authors derive pretraining recipes that achieve state-of-the-art multimodal performance using only 5% of the compute budget of naive joint training.

Technical Formulation

Let $\mathcal{D} = \{(x_i^L, x_i^V, x_i^G)\}$ be a dataset of language, visual understanding, and visual generation examples. A multimodal model f_θ is trained with a joint loss:

$$\mathcal{L}(\theta) = \alpha_L \mathcal{L}_L(\theta) + \alpha_V \mathcal{L}_V(\theta) + \alpha_G \mathcal{L}_G(\theta)$$

where $\alpha_L, \alpha_V, \alpha_G$ are modality weights. Standard approaches set these weights by data size. The paper proposes setting them adaptively based on the complexity and synergy of the current data mixture.

The synergy coefficient between modalities i and j at training step t is:

$$S_{ij}(t) = \frac{\partial \mathcal{L}_i / \partial \theta \cdot \partial \mathcal{L}_j / \partial \theta}{\|\partial \mathcal{L}_i / \partial \theta\| \cdot \|\partial \mathcal{L}_j / \partial \theta\|}$$

Positive S_{ij} indicates aligned gradients (synergy); negative indicates competition. Monitoring S_{ij} during training reveals when modalities are cooperating and when they are in conflict.

Learning or Inference Procedure

The study trains a series of controlled ablation models, holding constant the total compute budget while varying: (1) which modalities are present, (2) the order in which modalities are introduced, (3) the architecture of shared vs. modality-specific layers, and (4) the data mixture complexity. The final efficient recipe is validated by training 4×13.5B-parameter MoE models on 2T tokens.

What the Guarantee Says

No formal guarantee is provided. The empirical claim is that the four principles reliably improve multimodal model quality across the tested model sizes (up to 13.5B parameters) and data scales (up to 2T tokens). The cross-modal transfer measurements and vision laziness analysis provide mechanistic evidence for why the principles hold, but formal generalization bounds are absent.

Experimental Findings

- **Knowledge flow:** Pretraining on language improves visual understanding by up to 8 points (ImageNet accuracy) but improves visual generation by less than 2 FID points. Visual generation pretraining improves visual understanding by 3–5 points.
- **Synergy:** On complex datasets, joint training improves each modality by 3–12 points over single-modality training. On simple datasets, joint training hurts performance by 2–5 points.
- **Vision laziness:** Models with late visual integration have visual encoder representations that are 60–70% correlated with language encoder representations, vs. 30–40% for early-unified models — confirming that vision laziness reduces visual grounding.
- **Efficient recipes:** The derived recipe matches the quality of naive joint training at 5% compute, and exceeds sequential pretraining at all compute levels tested.

Ablations and Interpretation

Architecture: Fully shared architectures suffer modality competition even on complex data. Fully separate (modality-specific) architectures cannot share knowledge. The intermediate design — shared attention and normalization, modality-specific feed-forward — achieves the best trade-off in all settings.

Data mixing: The optimal mixture ratio is task-dependent, but the synergy coefficient S_{ij} predicts which mixtures will be beneficial in advance, enabling adaptive data scheduling without exhaustive grid search.

Training order: The vision laziness effect is most pronounced when vision is introduced after 30% of the total training budget has elapsed. Introducing vision at 10% or earlier eliminates the effect; full early unification (from step 0) is optimal.

Reference

Junlin Han, Shengbang Tong, David Fan, Minghao Chen, Philip Torr, Filippos Kokkinos, Mike Lewis. **Towards Physics of Multimodal Pretraining: Knowledge Flow, Modality Synergy, Early Unification, and Recipes**. arXiv:2608.05000, August 2026. https://arxiv.org/abs/2608.05000