

Keywords: LoRA • low-rank adaptation • Transformer attention • approximation theory

How Much Rank Does LoRA Need? Rank-Error Bounds for Transformer Attention

A Task-Dependent Theory Ties LoRA Rank to Downstream-Weighted Tail Energy — and Shows Softmax Saturation Can Shrink the Required Rank

By Gerard Conangla Planes

arXiv:2608.26052 • Preprint, August 2026

LoRA is the dominant paradigm for adapting large language models to downstream tasks, yet practitioners select the rank r by heuristic. This paper provides the first task-dependent theory of LoRA approximation error for Transformer attention: fixing a pretrained head, a target attention function, and a downstream input distribution, it bounds the minimum expected KL divergence achievable by a rank- r query LoRA update as an explicit function of the downstream-weighted tail energy of the target update matrix. A striking corollary is that softmax saturation can make the required rank strictly smaller than the rank needed to match the attention logits.

AT A GLANCE

Problem: LoRA rank r is chosen by trial-and-error despite being task-dependent; no theory predicts the KL error achievable at rank r .

Why it matters: Principled rank selection reduces parameter overhead while guaranteeing sufficient expressiveness for the adaptation.

Method: Bound the minimum expected KL error under rank- r query LoRA updates as a function of the downstream-weighted spectral tail energy T_r of the target update matrix.

Result: Softmax saturation lets rank needed to match the attention function be strictly lower than rank needed to match logits — constant separation for an explicit family.

Evidence: Theoretical bounds, extending to multi-head fused LoRA and joint query/key adaptation; construction proving tightness.

The Big Picture

LoRA (Low-Rank Adaptation) fine-tunes a pretrained weight matrix W_0 by adding $\Delta W = BA$ with $B \in \mathbb{R}^{d \times r}$ and $A \in \mathbb{R}^{r \times k}$, keeping W_0 frozen. The rank r controls both the parameter count and the class of reachable adaptations. In current practice, r is chosen from a small grid

$\{1, 2, 4, 8, 16\}$ by cross-validation, with no principled basis for the choice.

For Transformer attention, the query matrix W_Q determines the inner product structure of the attention mechanism. A rank- r update to W_Q shifts the key-query inner products in at most r directions. If the downstream task requires shifts in many independent directions, a small r is insufficient; if the task requires shifts in a low-dimensional subspace, a large r wastes parameters.

This paper asks: given the pretrained W_Q , the target W_Q^* , and the downstream input distribution $\mathcal{D}$, what is the minimum KL divergence between adapted and target attention distributions achievable by any rank- r query LoRA update?

Why Existing Approaches Fall Short

Existing analyses of LoRA operate at the level of weight-matrix approximation: a rank- r update approximates $\Delta W^* = W_Q^* - W_Q$ in spectral or Frobenius norm. But this misses two effects:

First, attention KL divergence depends on *downstream-weighted* approximation error. Components of ΔW^* that are large in matrix norm but project onto low-probability input directions have little effect on the adapted attention distribution. A matrix-norm-based analysis overestimates the required rank for these components.

Second, softmax applies a highly nonlinear map to the logits before computing the attention weights. Large logit differences saturate the softmax and make the attention distribution insensitive to further increases. A low-rank logit update can therefore achieve the same attention KL as a higher-rank update that matches the full logit change, because the relevant logit differences are already in the saturation regime.

Core Mechanism

The paper defines the **tail energy** T_r as the sum of eigenvalues beyond rank r in the downstream-weighted spectral decomposition of the target update:

$$T_r = \sum_{j>r} \lambda_j \left(\Sigma^{1/2} \Delta W^* \Delta W^* \Sigma^{1/2} \right)$$

where $\Sigma = \mathbb{E}_{\mathcal{D}}[xx^\top]$ is the downstream input covariance. The tail energy measures the fraction of the target adaptation that lies outside the subspace reachable by rank- r queries under the downstream distribution.

This quantity bridges weight-matrix approximation and attention KL divergence: a small T_r implies there exists a rank- r update with small expected KL error; a large T_r implies no rank- r update can achieve small error.

Technical Formulation

Fix an input x , query $q = W_Q x$, target query $q^* = W_Q^* x$, and key vectors $\{k_j\}$. The logit residual is $d = q - q^*$.

The KL divergence between the adapted and target attention distributions is:

$$\text{KL}(\text{softmax}(q^*) \parallel \text{softmax}(q^* - \Delta q)) = \psi(\|d\|_2)$$

where ψ is an increasing function bounded above by $\min\{\|d\|_2^2/4, \sqrt{2}\|d\|_2\}$.

For attention probabilities bounded away from zero (condition $\pi_j \geq \gamma > 0$ for all j), the lower bound is:

$$\text{KL} \geq \psi(\|d\|_2) \geq c_\gamma \|d\|_2$$

for an explicit constant c_γ depending on γ .

Under realizability, geometry, and moment conditions, the best rank- r expected KL error $\epsilon^*(r)$ satisfies:

$$c_1 \psi(\sqrt{T_r}) \leq \epsilon^*(r) \leq \min\left\{\frac{T_r}{4}, \sqrt{2T_r}\right\}$$

with explicit constants c_1 . These bounds are tight: the paper constructs families achieving both.

Learning or Inference Procedure

The theory yields a practical rank selection procedure:

1. Collect a sample $\{x_i\}$ from the downstream distribution
2. Estimate $\Sigma = \frac{1}{n} \sum_i x_i x_i^\top$
3. Compute the downstream-weighted SVD of the target update ΔW^* (estimated from a small fine-tuning probe if needed)
4. Evaluate T_r for each candidate rank r
5. Select r as the smallest rank with T_r below the desired KL error tolerance

At fine-tuning time, only the rank- r query LoRA is trained; all other weights remain frozen.

What the Guarantee Says

The bounds establish that the minimum KL error is controlled by T_r : if $T_r \leq \tau^2$, then a rank- r update can match the target attention to within $\min\{\tau^2/4, \sqrt{2}\tau\}$ expected KL. For saturated attention (large logit differences), the softmax nonlinearity allows this bound to be substantially tighter than the logit approximation error T_r alone would suggest.

The paper proves that there exist target functions where the attention-matching rank is a constant factor smaller than the logit-matching rank—closing a gap between empirical observations that low rank often suffices and the prior absence of theory explaining why.

Experimental Findings

The paper validates the rank-error bounds on two settings:

Synthetic attention heads: For planted low-rank targets, the measured KL error at each rank closely tracks $\min\{T_r/4, \sqrt{2T_r}\}$, confirming the upper bound is tight within a factor of 2.

LLM fine-tuning (RoBERTa-large, GLUE): The rank selected by the tail energy criterion achieves comparable task accuracy to a grid-searched rank in 8/9 GLUE tasks, while using 30–60% fewer LoRA parameters on average. On STS-B (a semantic similarity task where attention must capture fine-grained input directions), the criterion correctly predicts that rank 4 is insufficient and recommends rank 16.

Ablations and Interpretation

Softmax saturation effect: For heads with large pre-softmax logit magnitudes ($\|q^*\|_\infty > 10$), the attention-matching rank is consistently 2–4 $\times$ lower than the logit-matching rank, validating the theoretical separation.

Multi-head fused LoRA: When a single shared B is used across all heads, T_r must aggregate over head-specific downstream covariances; the bound remains tight under this extension.

Joint query/key updates: The analysis generalizes to adapting both W_Q and W_K simultaneously, with T_r defined over the joint key-query interaction tensor.

Distribution shift: If $\mathcal{D}$ is misspecified (e.g., using a generic corpus when the downstream distribution is specialized), T_r may underestimate the required rank. Domain-specific calibration data is therefore recommended for rank estimation.

Reference

Gerard Conangla Planes. **How Much Rank Does LoRA Need? Rank-Error Bounds for Transformer Attention.** arXiv:2608.26052, August 2026. <https://arxiv.org/abs/2608.26052>