Keywords: test-time reasoning • self-consistency • error localization • token efficiency

StepBack: Step-Level Error-Localized Resampling for Efficient Test-Time Reasoning

Localizing the First Error Step and Resampling Only Its Continuation Cuts Token Usage by Up to 41% at No Accuracy Cost

By Yang Ouyang, Jung-Eun Kim
Keywords: test-time reasoning • self-consistency • error localization • token efficiency

By Yang Ouyang, Jung-Eun Kim
0.6pt
Self-consistency generates many full chain-of-thought trajectories and takes a majority vote, but every chain runs to completion regardless of where the error occurs. StepBack localizes the most likely first error using lightweight signals and resamples only from that step onward, matching self-consistency accuracy on seven benchmarks while reducing token use by up to 41.26%.

AT A GLANCE

Problem: Self-consistency is expensive because every reasoning trajectory runs to completion, even when an error occurs early.
Why It Matters: Test-time compute budgets are a primary cost driver for deployed LLM reasoning; a method that preserves accuracy at lower token cost directly reduces inference cost.
Method: Localize the first likely error in draft trajectories using NLL and PRM signals, then resample only the erroneous step and its continuation, finishing with majority voting.
Result: Matches SC accuracy on 7 benchmarks; cuts token usage by up to 41.26%.
Evidence: GSM8K, MATH, AIME, MultiArith, HotpotQA, 2WikiMultiHopQA, MuSiQue; Llama-3.1-8B-Instruct backbone.

}{\NORMALSIZE\BFSERIES\COLOR{RULERED}}{}{0PT}{}%

[0.4pt]

}{0PT}{8PT}{4PT

0pt 4pt

By Yang Ouyang, Jung-Eun Kim

0.6pt

Self-consistency generates many full chain-of-thought trajectories and takes a majority vote, but every chain runs to completion regardless of where the error occurs. StepBack localizes the most likely first error using lightweight signals and resamples only from that step onward, matching self-consistency accuracy on seven benchmarks while reducing token use by up to 41.26%.

[backgroundcolor=boxgray,linecolor=rulered,linewidth=1pt, innerleftmargin=6pt,innerrightmargin=6pt, innertopmargin=6pt,innerbottommargin=6pt] **AT A GLANCE**[leftmargin=0pt,labelsep=4pt,itemsep=2pt,parsep=0pt, font=

THE BIG PICTURE

Chain-of-thought reasoning substantially improves LLM accuracy by externalizing intermediate steps. Self-consistency (SC) amplifies this further: by sampling many independent trajectories and taking a majority vote, SC cancels out reasoning errors that occur in different places on different runs. The cost is linear in the number of trajectories — a budget of 32 trajectories means 32

WHY EXISTING APPROACHES FALL SHORT

Self-consistency (Wang et al., 2022) achieves its accuracy benefit by diversity of reasoning paths, not by quality filtering. It treats every trajectory uniformly and cannot distinguish high-quality from low-quality prefixes.

Best-of-N and verifier-guided methods re-rank complete trajectories after generation. They still generate all N full chains and then select the best; the selection step does not reduce the generation cost.

Speculative decoding accelerates token generation but does not reduce the *number* of tokens in the reasoning chain, only the wall-clock time. It is orthogonal to StepBack.

Adaptive Consistency reduces the number of full trajectories by stopping early when the answer distribution is sufficiently peaked. StepBack and Adaptive Consistency are orthogonal: StepBack shortens each trajectory, while Adaptive Consistency reduces the trajectory count. The two compose.

CORE MECHANISM

Given a draft set of m trajectories, StepBack splits each trajectory using the model's natural delimiter tokens (newline, period, or explicit step markers).

For each step s_i^t, two signals are computed:

1. [itemsep=2pt,parsep=0pt]
2. **NLL signal:** The mean negative log-likelihood of tokens in s_i^t under the model, (s_i^t).
3. **PRM signal:** The process reward model score (s_i^t), estimating the probability that all steps up to and including s_i^t are correct.

The anchor step s_i^a_i is the earliest step satisfying: \[(s_i^a_i) > \;\;\;\; (s_i^a_i) < , \] with = 0.1 and = 0.05. If no step is flagged, the trajectory is accepted as-is.

For each flagged trajectory, k = 2 fresh continuations are sampled from the anchor step onward, keeping the prefix s_i^ fixed. The final answers from all trajectories (original and resampled) are pooled and the plurality answer is returned.

TECHNICAL FORMULATION

Let = \ be the initial set of m = 16 draft trajectories. For trajectory and the original continuation is c_i = s_i^a_i _i^j (\ (p_i _i^j) : ,\, j=1,. \] The returned answer is [a' = a].

The expected token count per input is: \[[] = m _ + _, \] where _, _ are mean prefix and suffix lengths. When errors cluster early (_ small), the saving is largest.

LEARNING OR INFERENCE PROCEDURE

StepBack is a pure inference-time method; no model parameters are modified. The required inputs are:

The PRM used in experiments is a Llama-3.1-8B model fine-tuned on step-level correctness labels from the PRM800K dataset.

WHAT THE GUARANTEE SAYS

No formal guarantee is claimed. The empirical claim is: on the seven benchmarks tested, StepBack's accuracy is not statistically significantly different from SC at the same sampling budget of m = 32 effective trajectories. The token-reduction claim (up to 41.26%) is an empirical upper bound observed on the MATH benchmark; the reduction is smaller on benchmarks where errors occur later in trajectories.

EXPERIMENTAL FINDINGS

ABLATIONS AND INTERPRETATION

NLL-only vs.\ PRM-only vs.\ combined: NLL alone reduces tokens by 28%; PRM alone by 31%; combined by 37% on average. The two signals flag different steps on roughly 20% of trajectories, confirming complementarity.

Resampling k: k = 1 achieves 90% of the token savings at a 1–2% accuracy drop; k = 3 shows no further accuracy gain over k = 2 while increasing cost.

Step granularity: Fine-grained steps (sentence level) outperform coarse steps (paragraph level) by 1.3% average accuracy; the localized error is more precisely identified at finer granularity.

PRM quality: Replacing the trained PRM with a zero-shot LLM self-evaluation of step correctness reduces accuracy by 2.1% and token savings by 8%, confirming the importance of a dedicated PRM.

{REFERENCE

Yang Ouyang, Jung-Eun Kim. **StepBack: Step-Level Error-Localized Resampling for Efficient Test-Time Reasoning.** arXiv:2611.08403. NeurIPS 2026.
https://arxiv.org/abs/2611.08403